

Car CO2 Emission Prediction (Regression)

About this notebook

In this notebook, I will use data about cars and regression models to predict CO2 emission. This data provides model-specific fuel consumption ratings and estimated carbon dioxide emissions for new light-duty vehicles for retail sale in Canada. Note: This is a part of IBM Data Scientist course projects.

Table of contents

1. Variable Descriptions2. Understand Data3. Linear Regression4. Non-linear Regression4.1 Polynomial Regression4.2 Ridge Regression

</div>

```
In [2]: #import packages
import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
%matplotlib inline
import seaborn as sns
```

```
In [3]: #get data
df=pd.read_csv('https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/M
L0101ENV3/labs/FuelConsumptionCo2.csv')
```

This dataset has 1,067 observations and 13 columns.

Variable Descriptions

We have downloaded a fuel consumption dataset, **FuelConsumption.csv** , which contains model-specific fuel consumption ratings and estimated carbon dioxide emissions for new light-duty vehicles for retail sale in Canada. [Dataset source](http://open.canada.ca/data/en/dataset/98f1a129-f628-4ce4-b24d-6f16bf24dd64) (<http://open.canada.ca/data/en/dataset/98f1a129-f628-4ce4-b24d-6f16bf24dd64>).

- **MODELYEAR** e.g. 2014
- **MAKE** e.g. Acura
- **MODEL** e.g. ILX
- **VEHICLE CLASS** e.g. SUV
- **ENGINE SIZE** e.g. 4.7
- **CYLINDERS** e.g 6
- **TRANSMISSION** e.g. A6
- **FUEL CONSUMPTION in CITY(L/100 km)** e.g. 9.9
- **FUEL CONSUMPTION in HWY (L/100 km)** e.g. 8.9
- **FUEL CONSUMPTION COMB (L/100 km)** e.g. 9.2
- **CO2 EMISSIONS (g/km)** e.g. 182

Understand Data

In [12]:

```
#shape of the data
df.shape
```

Out[12]: (1067, 13)

In [10]:

```
#the head of data
df.head()
```

Out[10]:

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUELTYPE	FUELCONSUMPTION_CITY
0	2014	ACURA	ILX	COMPACT	2.0	4	AS5	Z	13.296532
1	2014	ACURA	ILX	COMPACT	2.4	4	M6	Z	13.296532
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7	Z	10.250000
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6	Z	15.550000
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS6	Z	15.550000

In [103]:

```
df.dtypes
```

Out[103]:

MODELYEAR	int64
MAKE	object
MODEL	object
VEHICLECLASS	object
ENGINE SIZE	float64
CYLINDERS	int64
TRANSMISSION	object
FUELTYPE	object
FUELCONSUMPTION_CITY	float64
FUELCONSUMPTION_HWY	float64
FUELCONSUMPTION_COMB	float64
FUELCONSUMPTION_COMB_MPG	int64
CO2EMISSIONS	int64
dtype:	object


In [13]:

```
#see the descriptive statistics of the numeric variables
df.describe()
```

Out[13]:

	MODELYEAR	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY	FUELCONSUMPTION_COMB
count	1067.0	1067.000000	1067.000000	1067.000000	1067.000000	1067.000000
mean	2014.0	3.346298	5.794752	13.296532	9.474602	11.000000
std	0.0	1.415895	1.797447	4.101253	2.794510	3.000000
min	2014.0	1.000000	3.000000	4.600000	4.900000	4.000000
25%	2014.0	2.000000	4.000000	10.250000	7.500000	9.000000
50%	2014.0	3.400000	6.000000	12.600000	8.800000	10.000000
75%	2014.0	4.300000	8.000000	15.550000	10.850000	13.000000
max	2014.0	8.400000	12.000000	30.200000	20.500000	25.000000


```
In [19]: #Generate scatter plots for numeric variables VS CO2EMISSIONS to see if correlations
sns.pairplot(df, x_vars=["ENGINE_SIZE", "CYLINDERS", "FUELCONSUMPTION_COMB"], y_vars=["CO2EMISSIONS"],
            height=5, aspect=.8, kind="reg");
```



CYLINDERS are categorical. There are 5 distinct values in the dataset. We want to see if there are any difference in the mean of each type of cylinder numbers realting to CO2EMISSIONS.

```
In [25]: #box plot for CYLINERS
sns.boxplot(x="CYLINDERS", y="CO2EMISSIONS", data=df)
```

```
Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x20c70bbf630>
```



So, the number of cylinders has some influence in CO2EMISSIONS.

There are four columns about fuel consumptions: FUELCONSUMPTION_CITY, FUELCONSUMPTION_HWY, FUELCONSUMPTION_COMB, and FUELCONSUMPTION_COMB_MPG. If they are highly correlated, we can eliminate some of them to make the model lighter. We will run a correlation.

```
In [29]: df.corr()
```

```
Out[29]:
```

	MODELYEAR	ENGINESIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_
MODELYEAR	NaN	NaN	NaN	NaN	
ENGINESIZE	NaN	1.000000	0.934011	0.832225	0.77
CYLINDERS	NaN	0.934011	1.000000	0.796473	0.72
FUELCONSUMPTION_CITY	NaN	0.832225	0.796473	1.000000	0.96
FUELCONSUMPTION_HWY	NaN	0.778746	0.724594	0.965718	1.00
FUELCONSUMPTION_COMB	NaN	0.819482	0.776788	0.995542	0.98
FUELCONSUMPTION_COMB_MPG	NaN	-0.808554	-0.770430	-0.935613	-0.85
CO2EMISSIONS	NaN	0.874154	0.849685	0.898039	0.86

So, we can use one of the columns about fuel consumption without lossing too much information.Now we know ENGINESIZE, CYLINERS, and FUELCOMSUMPTION_COMB are useful for building the regression model.

Linear Regression


```
In [4]: #create a new dataframe for model building
lrdf = df[['ENGINESIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS']]
lrdf.shape
```

```
Out[4]: (1067, 4)
```

```
In [5]: #spilt train and test data
msk = np.random.rand(len(df)) < 0.8
train = lrdf[msk]
test = lrdf[~msk]
```

```
In [37]: #train data distribution
sns.pairplot(lrdf, x_vars=['ENGINESIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB'], y_vars=['CO2EMISSIONS'],
             height=5, aspect=.8, kind="reg")
```

```
Out[37]: <seaborn.axisgrid.PairGrid at 0x20c70c95f60>
```



```
In [6]: #build the model
#step 1, import packages
from sklearn.linear_model import LinearRegression
lm=LinearRegression()

#step 2, create independent variables
X=train[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']]
Y=train[['CO2EMISSIONS']]
```

```
In [47]: #step 3, fit the linear regression model
lm.fit(X,Y)
# The coefficients
print ('Coefficients: ', lm.coef_)
print ('Intercept: ',lm.intercept_)
```

```
Coefficients: [[ 9.92107897  7.75011521 10.0252928 ]]
Intercept: [61.92097822]
```

Now the model equation is: $Y = 61.92097822 + 9.92107897ENGINE_SIZE + 7.75011521CYLINDERS + 10.0252928 \cdot FUELCONSUMPTION_COMB + \text{Error}$

```
In [48]: #model evaluation
from sklearn.metrics import r2_score

test_x = test[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']]
test_y = test[['CO2EMISSIONS']]
test_y_hat = lm.predict(test_x)

print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_hat - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_hat - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_hat , test_y) )
```

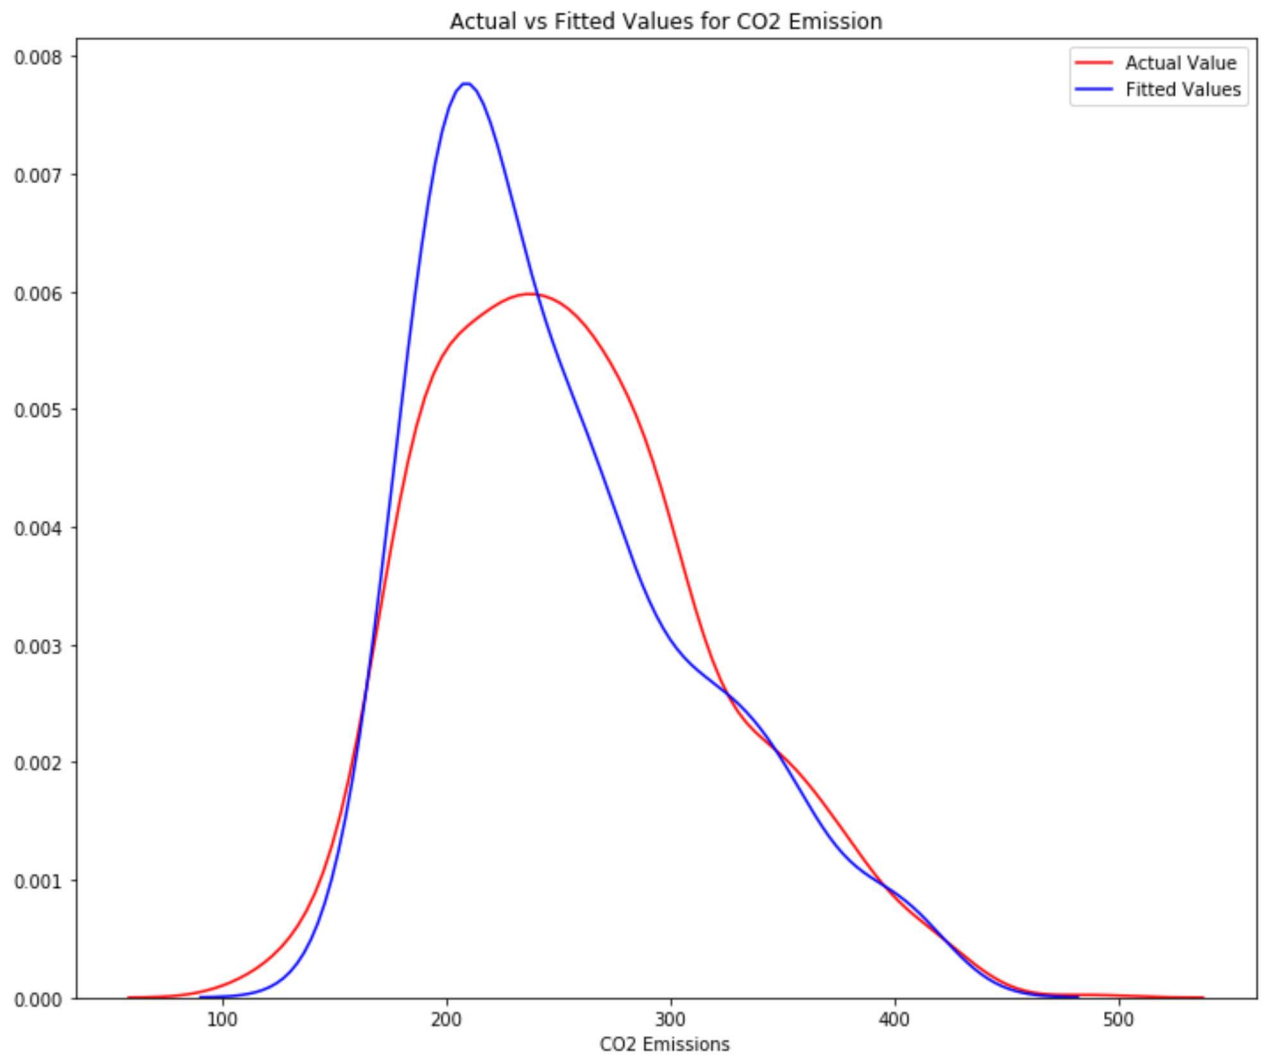
```
Mean absolute error: 16.67
Residual sum of squares (MSE): 527.73
R2-score: 0.85
```

```
In [51]: #Distribution Plot
width = 12
height = 10
plt.figure(figsize=(width, height))

ax1 = sns.distplot(lrdf['CO2EMISSIONS'], hist=False, color="r", label="Actual Value")
sns.distplot(test_y_hat, hist=False, color="b", label="Fitted Values" , ax=ax1)

plt.title('Actual vs Fitted Values for CO2 Emission')
plt.xlabel('CO2 Emissions')
plt.ylabel('')

plt.show()
plt.close()
```



Non Linear Regression

Although the above linear model fits well (R^2 is 0.85), as seen from the distribution plot, it fails to explain some data points. This time we try non-linear regression models.

Polynomial Model

```
In [7]: #import packages
from sklearn.preprocessing import PolynomialFeatures
from sklearn import linear_model
```

```
In [8]: #first, try ENGINE SIZE as the only independent variable
train_x=train[['ENGINE SIZE']]
train_y=train[['CO2 EMISSIONS']]
poly = PolynomialFeatures(degree=2)
train_x_poly = poly.fit_transform(train_x)
train_x_poly
```


```
Out[8]: array([[ 1. ,  2.4 ,  5.76],
 [ 1. ,  1.5 ,  2.25],
 [ 1. ,  3.5 , 12.25],
 ...,
 [ 1. ,  3.2 , 10.24],
 [ 1. ,  3.2 , 10.24],
 [ 1. ,  3.2 , 10.24]])
```

```
In [59]: train_y_ = lm.fit(train_x_poly, train_y)
# The coefficients
print ('Coefficients: ', lm.coef_)
print ('Intercept: ',lm.intercept_)
```

```
Coefficients: [[ 0.          51.17136746 -1.63000129]]
Intercept: [105.69276866]
```

```
In [60]: plt.scatter(train.ENGINE SIZE, train.CO2 EMISSIONS, color='blue')
XX = np.arange(0.0, 10.0, 0.1)
yy = lm.intercept_[0]+ lm.coef_[0][1]*XX+ lm.coef_[0][2]*np.power(XX, 2)
plt.plot(XX, yy, '-r' )
plt.xlabel("Engine size")
plt.ylabel("Emission")
```

```
Out[60]: Text(0, 0.5, 'Emission')
```



```
In [62]: #Evaluation
test_x = test[['ENGINE SIZE']]
test_y = test[['CO2 EMISSIONS']]
test_x_poly = poly.fit_transform(test_x)
test_y_ = lm.predict(test_x_poly)

print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

```
Mean absolute error: 20.65
Residual sum of squares (MSE): 737.28
R2-score: 0.74
```

```

In [13]: #Optimize the polynomial degrees
Rsqu_test = []

order = [1, 2, 3, 4, 5, 6, 7]
for n in order:
    poly = PolynomialFeatures(degree=n)
    train_x=train[['ENGINE SIZE']]
    train_y=train[['CO2EMISSIONS']]
    train_x_poly = poly.fit_transform(train_x)
    test_x = test[['ENGINE SIZE']]
    test_y = test[['CO2EMISSIONS']]
    test_x_poly = poly.fit_transform(test_x)
    #Lr=LinearRegression()
    lm.fit(train_x_poly, train_y)
    Rsqu_test.append(lm.score(test_x_poly, test_y))

plt.plot(order, Rsqu_test)
#plt.figure(order, Rsqu_test, figsize=(3,4))
plt.xlabel('order')
plt.ylabel('R^2')
plt.title('R^2 Using Test Data')

```

Out[13]: Text(0.5, 1.0, 'R^2 Using Test Data')



When the polynomial degree for ENGINE SIZE is 4, we have a good R^2 . Although it the R^2 increases when the degree increases, we use degree=4 because it's not necessary to let the degree equal to 7 or even higher. The model may fit the noise too well. Let's have a look of the distribution plot.


```
In [81]: #degree equals to 4
poly = PolynomialFeatures(degree=4)
train_x=train[['ENGINE SIZE']]
train_y=train[['CO2EMISSIONS']]
train_x_poly = poly.fit_transform(train_x)
test_x = test[['ENGINE SIZE']]
test_y = test[['CO2EMISSIONS']]
test_x_poly = poly.fit_transform(test_x)
lm.fit(train_x_poly, train_y)
test_y_hat=lm.predict(test_x_poly)

width = 12
height = 10

plt.figure(figsize=(width, height))


ax1 = sns.distplot(lrd['CO2EMISSIONS'], hist=False, color="r", label="Actual Value")
sns.distplot(test_y_hat, hist=False, color="b", label="Fitted Values" , ax=ax1)

plt.title('Actual vs Fitted Values for CO2 Emission')
plt.xlabel('CO2 Emissions')
plt.ylabel('')

plt.show()
plt.close()
```

D:\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



Ridge Regression

We will use Ridge Regression to improve the model. Ridge Regression controls the degree of polynomial regression by the parameter Alpha. Ridge Regression has the following benefits: -reduce the multicollinearity of endogenous variables in models -reduce model complexity and prevent over-fitting which may result from simple linear regression.

```
In [123]: #prepare data
x_data = lrd[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']]
y_data= lrd[['CO2EMISSIONS']]
train_x, test_x, train_y, test_y = train_test_split(x_data, y_data, test_size=0.2, random_state=0)
poly=PolynomialFeatures(degree=2)
lm=LinearRegression()
```

```
In [117]: #import packages
from sklearn.linear_model import Ridge
```

Let's create a Ridge regression object, setting the regularization parameter to 0.1

```
In [118]: RigeModel=Ridge(alpha=0.1)
```

```
In [124]: train_x_poly= poly.fit_transform(train_x)
test_x_poly = poly.fit_transform(test_x)
```

```
In [125]: RigeModel.fit(train_x_poly, train_y)
yhat = RigeModel.predict(test_x_poly)
```

```
In [127]: print('predicted:', yhat[0:4])
print('test set :', test_y[0:4].values)
```


```
predicted: [[346.34060871]
 [211.77465377]
 [224.66033349]
 [221.681592  ]]
test set : [[356]
 [209]
 [230]
 [212]]
```

```
In [128]: Rsqu_test = []
Rsqu_train = []
dummy1 = []
ALFA = 10 * np.array(range(0,1000))
for alfa in ALFA:
    RigeModel = Ridge(alpha=alfa)
    RigeModel.fit(train_x_poly, train_y)
    Rsqu_test.append(RigeModel.score(test_x_poly, test_y))
    Rsqu_train.append(RigeModel.score(train_x_poly, train_y))
```

```
In [129]: #plot for different alfa
width = 12
height = 10
plt.figure(figsize=(width, height))

plt.plot(ALFA,Rsqu_test, label='validation data ')
plt.plot(ALFA,Rsqu_train, 'r', label='training Data ')
plt.xlabel('alpha')
plt.ylabel('R^2')
plt.legend()
```

Out[129]: <matplotlib.legend.Legend at 0x20c72115c18>



The blue line represents the R^2 of the test data, and the red line represents the R^2 of the training data. The x-axis represents the different values of Alpha. The red line in represents the R^2 of the test data, as Alpha increases the R^2 decreases; therefore as Alpha increases the model performs worse on the test data. The blue line represents the R^2 on the validation data, as the value for Alpha increases the R^2 decreases.

Lower alpha value (lower polynomial degree) will yield better R^2 . Let's find the exact alpha value and then determine the model.

```
In [189]: x_data = lrdff[['ENGINE SIZE', 'CYLINDERS', 'FUEL CONSUMPTION COMB']]
y_data= lrdff[['CO2 EMISSIONS']]
train_x, test_x, train_y, test_y = train_test_split(x_data, y_data, test_size=0.2, random_state=0)
```

```
In [ ]: alphas = 10**np.linspace(10,-2,100)*0.5
alphas
```

```
In [190]: ridge = Ridge(normalize = True)
          coefs = []

          for a in alphas:
              ridge.set_params(alpha = a)
              ridge.fit(train_x, train_y)
              coefs.append(ridge.coef_)

          np.shape(coefs)
```

Out[190]: (100, 1, 3)

```
In [181]: from sklearn.preprocessing import scale
          from sklearn.model_selection import train_test_split
          from sklearn.linear_model import Ridge, RidgeCV, Lasso, LassoCV
          from sklearn.metrics import mean_squared_error
```

```
In [191]: #obtain the best alpha by cross-validation
          ridgecv = RidgeCV(alphas = alphas, scoring = 'neg_mean_squared_error', normalize = True)
          ridgecv.fit(train_x, train_y)
          ridgecv.alpha_
```

Out[191]: 0.01155064850041579

Therefore, we see that the value of alpha that results in the smallest cross-validation error is 0.012.

```
In [192]: #obtain the RMSE
          ridge = Ridge(alpha = ridgecv.alpha_, normalize = True)
          ridge.fit(train_x, train_y)
          mean_squared_error(test_y, ridge.predict(test_x))
```

Out[192]: 593.7935827354731

```
In [195]: #Use the optimized alpha for the new ridge regression model using all data
          ridge_final = Ridge(alpha=0.012)
          ridge_final.fit(x_data, y_data)
          ridge_final.coef_
```

Out[195]: array([[10.85494779, 7.51636562, 9.5956638]])

Is it necessary to use Ridge Regression than Linear Regression? We compare the MSE. The MSE obtained from the linear regression model is 737.28, however MSE from Ridge Regression is much better: 593.79. So in this case, we use the Ridge Regression model.